

Blockchains and State Machines

An implementation of the PinkScorpion protocol

Gary Mawdsley CTO/CEO Lockular Limited

May 8th 2024¹

¹ Work in progress

This paper outlines the PinkScorpion protocol implementation using a blockchain based state machine to model the transitions defined by the protocol, leading to an immutable log of file activity.

Blockchains are state machines where state transitions are recorded on an immutable ledger. The transitions can be complex and governed by a set of rules including validation of actors undertaking the transitions. A transition is also known as a transaction. The ledger is not only immutable but is replicated and many nodes work to form a single blockchain.

The State Machine

A state machine is crucial for maintaining the integrity and consistency of the ledger across all nodes in the network. It is relevant in:

- **State Management:** A blockchain is a state machine where transactions transition the system from one state to another. Each block represents a series of state changes.
- **Consensus:** Blockchains use consensus algorithms (like Proof of Work, Proof of Stake) to agree on the next valid block (and thus the next state of the ledger). This ensures all nodes maintain a consistent view of the state.
- **Smart Contracts:** On platforms like Polkadot, smart contracts are programs that run on the blockchain's decentralised state machine. They execute and manage state transitions based on the code's logic.
- **Immutability and Security:** The state machine approach helps in ensuring that once a state transition (like a transaction) is made and agreed upon, it cannot be altered, providing a secure and immutable ledger.
- **Scalability and Efficiency:** Optimisations in the state transition process (like sharding or state channels) can enhance the scalability and efficiency of blockchain operations.

In summary, the state machine concept is integral to how blockchains operate, ensuring secure, consistent, and decentralized management

of digital assets and data. Smart Contracts afford a means to influence the subject of a specific state machine but may not offer the efficiencies of writing a bespoke blockchain.

Scalability and Efficiency

Sharding

Sharding is a method that partitions a blockchain into smaller, more manageable pieces called "shards." Each shard contains its own independent state and transaction history, which allows multiple transactions to be processed in parallel, rather than sequentially. This parallel processing capability significantly increases the throughput of the blockchain.

- How does it work: The blockchain network is divided into several shards, with each shard handling a portion of the network's transactions and state. Nodes in the network only need to process and store information related to their respective shard rather than the entire network.
- Benefits: Sharding drastically improves the scalability of a blockchain network because it reduces the workload on individual nodes, which leads to faster transaction processing times and lower costs.
- Challenges: Sharding does introduce complexity in managing cross-shard communication and ensuring security across shards, as the smaller number of nodes in each shard can potentially make them more susceptible to attacks. Adopting a platform where sharding is intrinsic mitigates much of the risk.

State Channels

State channels are a second-layer scaling technique that allows participants to conduct transactions off the main blockchain (off-chain) in a private channel, thus freeing up the network's resources. Transactions within a state channel do not need to be broadcast to the entire network and are only settled on the blockchain after the channel is closed. This is very similar to how Lockular's filesystem change log is cast into the blockchain where groups of file operations are batched, hashed and recorded as a single blockchain transaction.

- How does it work: Two or more parties open a state channel by committing a transaction to the ledger indicating the opening of the channel. The parties can then perform an unlimited number of transactions amongst themselves, which are instant and free from

transaction fees. Once all transactions are complete, the final state of these transactions is committed to the blockchain.

- **Benefits:** State channels greatly reduce the burden on the blockchain, as only two transactions (open and close) are recorded on-chain. Transactions within the channel are extremely fast and do not incur direct costs per transaction, making them ideal for scenarios requiring high transaction volumes or rapid interactions.
- **Challenges:** The main challenges with state channels include the need for all parties to remain online to sign off on transactions and potential issues in disputes or if one party tries to close the channel unilaterally with incorrect data.

Both sharding and state channels provide effective strategies to enhance the efficiency and scalability of blockchains, making them capable of handling higher transaction volumes and providing faster processing times while maintaining security and decentralisation.

Polkadot and Sharding

Parachains in the Polkadot ecosystem are analogous to sharding as defined in a general blockchain architecture. Parachains are individual blockchains that run in parallel within the Polkadot network, each with its own unique features and purposes but connected and secured by the Polkadot Relay Chain. How Parachains Work:

- **Parallel Processing:** Each parachain processes transactions independently, allowing for parallel data processing. This is similar to how sharding works, where different shards handle different transactions simultaneously.
- **Shared Security:** Parachains benefit from the shared security model of the Polkadot network. The Relay Chain provides security to all connected parachains, which means individual parachains do not need to provide their own security measures.
- **Interoperability:** Parachains are inherently designed to be interoperable within the Polkadot ecosystem. They can communicate and transfer data or assets seamlessly through the Polkadot Relay Chain, using cross-chain message passing (XCMP).

Benefits of Parachains:

- **Scalability:** By processing transactions on different parachains simultaneously, Polkadot can handle a much higher throughput than traditional single-chain architectures.

- Specialisation: Each parachain can be designed for specific use cases with optimised functionality, whether for finance, identity management, data storage, etc.
- Flexibility: Developers can build parachains with different governance structures or features that best suit their needs, without conforming to the constraints of a one-size-fits-all model.

Challenges:

- Complexity: Managing a network of multiple parachains adds layers of complexity in terms of governance, communication, and operation.
- Resource Intensive: Launching and maintaining a parachain can be resource-intensive, as it requires securing a slot through Polkadot's parachain slot auctions, which can be competitive and costly.

The CoreTime architecture will transform the ease of parachain validation.

In summary, parachains are Polkadot's approach to achieving a scalable, interoperable, and secure multi-chain ecosystem, similar to how sharding aims to scale blockchain networks by dividing them into more manageable parts.

Polkadot and State Channels

Polkadot itself does not natively support state channels directly within its core protocol. However, the Polkadot ecosystem is designed to be highly flexible and extensible, allowing for various layer-two solutions, including state channels, to be implemented on top of or alongside it through parachains or external networks. Here's how one could implement State Channels in Polkadot:

- Parachains: Developers can create specific parachains that implement state channel technology. These parachains can handle the off-chain transactions and only interact with the main Polkadot Relay Chain for final settlements or dispute resolutions.
- External Solutions: Projects could also integrate existing state channel solutions that are not natively part of the Polkadot ecosystem but can interact with it through cross-chain communication. This could involve using bridges or other interoperability protocols to connect Polkadot with networks that support state channels.

We know Aleph is a fast blockchain integrated by a bridge. This may be a route to textitstate channels where desirable.

Potential for Future Developments:

- Substrate: Since Polkadot is built using Substrate, a blockchain-building framework, we have the tools to create custom parachains with built-in support for state channels. This flexibility allows

Lockular the potential development and integration of state channels helping to grow and evolve the ecosystem.

This would be a bespoke build as opposed to leveraging something say Aleph have already?

In Summary, whilst Polkadot does not currently have built-in support for state channels, its architecture allows for the integration of such technologies through parachains or external solutions, providing a pathway for leveraging state channels within the Polkadot ecosystem. Also note, in general state channels hold a major requirement for deposits to be in place in order to be really secure. With Polkadot it is perfectly possible to build state channels on top of parachains. State channels in general also have limitations in terms of extensibility as a whole where Polkadot does not, since it can theoretically support any state transition logic.

Bespoke State Machines

As discussed the fundamental principle of a blockchain is to control and immutably account for state transitions in a state machine. Here are a few scenarios where designing a bespoke state machine and controlling its transitions could be beneficial:

- Supply Chain Management: Creating a blockchain system to track and validate the movement of goods (represented as NFTs) in a supply chain. Customising state transitions enables verification and updating of the product's state (shipped, received, etc.) as it moves through different stages.
- Identity and Access Control: Designing a decentralised identity management system where state transitions occur when users update their information, grant access, or revoke permissions.
- Secure Design Services: Tracking the collaboration of low level design assets and their aggregation into fuller design outputs, e.g. Silicon Chip design.
- Gaming and NFTs: Designing a state machine for a gaming platform where ownership, trading, or evolution of in-game assets (NFTs) is managed on the blockchain. Defining the state transitions ensures fairness and integrity within the gaming ecosystem.

This sector is Lockular's primary provenance focus

Lockular's Workflow Actor Management component adds a smart contract solution here

Lockular's provenance tracking filesystem tracks low level design from conception to output

In these cases, designing a bespoke state machine allows logic, rules, and security measures to be tailored to suit specific use cases, providing more control and customisation over how the system evolves over time.

Polkadot and bespoke blockchain validation

Bespoke blockchains are intrinsically supported by Polkadot Substrate. Validators in a blockchain network verify and validate transactions and state changes based on the predefined rules and logic within the state machine. For bespoke state machines like those in Substrate-based blockchains, validators follow a generic validation process that includes a few key steps:

- **Consensus Rules:** Validators ensure that all transactions and state transitions comply with the consensus rules established by the network. These rules define the validity of blocks and transactions, ensuring agreement among nodes.
- **Runtime Logic:** Validators run the custom logic encoded in the runtime modules of the state machine. This logic determines the validity of state transitions, ensuring that changes adhere to the predefined rules set by the blockchain's design.
- **Signature Verification:** Validators check the cryptographic signatures of transactions and blocks to confirm that they are authorised and originate from valid sources, preventing unauthorised modifications.
- **Finality and Agreement:** Validators participate in the consensus mechanism of the network to agree on the state changes and ensure that all honest nodes converge to the same state.
- **Execution of State Transitions:** Validators execute and apply state transitions according to the rules defined in the state machine. They verify that the changes are consistent, valid, and do not violate any protocol rules.

In summary, validators play a crucial role in validating state changes (transactions) within bespoke state machines by enforcing consensus, executing runtime logic, verifying signatures, and ensuring the integrity and security of the blockchain network.

Designing a blockchain

When designing a blockchain using the Substrate framework, there's the opportunity to define not just the structure of the blockchain itself but also the rules and logic governing its state changes. This includes creating a bespoke state machine tailored to the specific needs of the use case.

These are the high level steps:

A smart contract can often be used instead of a bespoke blockchain. However this may limit scalability option viz a v sharding and state channels

- **Define Blockchain Structure:** Start by setting up the fundamental structure of the blockchain using Substrate. This involves configuring the consensus mechanism, defining the roles of participants (validators, nominators, etc.), and establishing the base functionality.
- **Customise State Transition Logic:** Using Substrate's runtime modules, define the logic for state transitions. This involves specifying how the data on the blockchain can change over time, implementing rules for transactions, managing accounts, defining governance processes, or any other functionality specific to our use case.
- **Security and Validation:** As part of this process, ensure that the state transitions are secure, follow the predefined rules, and are validated by the network's validators to maintain consistency and integrity.
- **Runtime Upgradability:** Substrate also offers the advantage of runtime upgradability, allowing one to modify or upgrade the state machine logic without hard forks, providing flexibility and adaptability as our application evolves.

By defining a bespoke blockchain and its bespoke state machine, one has the autonomy to create a blockchain solution that precisely fits the use case requirements, whether it's a decentralised application, a specialised financial system, supply chain management, or any other use case you envision.

Hooking in to validators

Presently in the Polkadot ecosystem, parachains (like our bespoke blockchain) secure slots on the Polkadot relay chain through a mechanism called parachain slot auctions. Once our parachain secures a slot on the relay chain, it gains the ability to interact with and benefit from the security provided by the relay chain validators.

When our parachain is connected to the Polkadot relay chain then:

- **Validation through Relay Chain Validators:** The relay chain validators are responsible for validating and securing the entire Polkadot network, including parachains. Once our parachain is connected, the validators of the relay chain will also validate the transactions and state changes within our parachain.
- **Security of the Parachain:** our parachain benefits from the security and consensus provided by the relay chain validators. They ensure the correctness and integrity of state transitions within our parachain, adding to the overall security of our blockchain.

- **Cross-Chain Communication:** Being part of the Polkadot network allows our parachain to communicate and transact with other parachains and the relay chain, enabling interoperability and collaboration between different blockchains.

Securing a slot on the relay chain grants a parachain access to the validation and security infrastructure provided by the Polkadot network, allowing a bespoke blockchain to operate securely within the larger ecosystem.

Securing a Parachain slot at auction can be expensive and lumpy. Coretime was announced in 2023 allowing access to validators via highly granular validation time slots offered as NFTs

But precisely what do validators validate?

Blocks are formed by collators; collators are essentially nodes that are operated by the developers or network participants of the bespoke parachain (built using Substrate or another compatible framework). These nodes are responsible for producing blocks, maintaining the state of the parachain, and interacting with the Polkadot relay chain.

When building parachains using Substrate one typically writes code that defines the behaviour of collators. This code handles the following tasks:

- **Block Production:** Collecting transactions, creating new blocks, and proposing them to the relay chain for inclusion.
- **State Maintenance:** Executing the logic defined in the runtime modules to process transactions and update the state of the parachain.
- **Interaction with Relay Chain:** Communicating with the Polkadot relay chain by submitting proposed blocks and handling interactions between the parachain and the relay chain validators.

Collators are an integral part of a parachain's infrastructure. Let's look at a real example albeit a high level.

Example bespoke blockchain in Substrate

Below is a simplified example of how a Substrate-based blockchain could be structured to include the logic for collators. This example assumes the creation of a simple custom parachain using Substrate's FRAME framework:

```

1 // Define the bespoke parachain's runtime modules
2 pub mod my_parachain {
3     use frame_system::Config;
4     use frame_support::dispatch;
5
6     // Define the configuration for our runtime
7     pub trait Config: frame_system::Config {}
8
9     // Define our parachain's custom module

```

but where in the example is the stuff required to be run by the collators - point it out


```

10     pub mod custom_module {
11         use super::*;
12         use frame_support::{decl_module, decl_storage, dispatch::
DispatchResult};
13
14         // Define our module's storage items
15         pub trait Config: my_parachain::Config {}
16         decl_storage! {
17             trait Store for Module<T: Config> as MyParachainModule {
18                 // Define our storage items here
19                 ExampleValue: u32;
20             }
21         }
22         // Define our module's dispatchable functions
23         decl_module! {
24             pub struct Module<T: Config> for enum Call where origin: T::
Origin {
25                 // Initialize the module
26                 fn deposit_event() = default;
27
28                 // Example function that updates storage
29                 #[weight = 10_000]
30                 pub fn update_value(origin, new_value: u32) -> dispatch
::DispatchResult {
31                     let _sender = ensure_signed(origin)?;
32                     <ExampleValue<T>::put(new_value);
33                     Ok(())
34                 }
35             }
36         }
37     }
38 }
39
40 // Define the parachain runtime by composing the modules
41 parameter_types! {
42     pub const BlockPeriod: u64 = 6;
43 }
44
45 impl my_parachain::Config for Runtime {
46     // Implement our parachain-specific configuration
47     // ...
48 }
49
50 impl pallet_timestamp::Config for Runtime {
51     type MinimumPeriod = BlockPeriod;
52     // ...
53 }
54
55 // Complete the parachain runtime setup and construct it
56 construct_runtime!(
57     pub enum Runtime where
58         Block = Block,
59         NodeBlock = opaque::Block,
60         UncheckedExtrinsic = UncheckedExtrinsic
61     {
62         // Add FRAME modules and our custom parachain modules
63         System: frame_system::{Module, Call, Config},
64         Timestamp: pallet_timestamp::{Module, Call, Storage, Config<T>},
65         MyParachainModule: my_parachain::{Module, Call, Storage, Config<
T>},
66     }

```

```

67 );
68
69 // Implement collator-specific logic (block production, communication
70 // with relay chain, etc.)
71 // ...
72

```

Listing 1: FRAME framework example

This is a basic outline demonstrating the structure of a Substrate-based parachain. In practice, one needs to define the collator-specific logic within the parachain runtime, specifying how blocks are produced, the communication protocol with the relay chain, and other functionalities related to collators' roles. The `myparachain` module is a placeholder for the required custom parachain logic. Within this module storage items, dispatchable functions, and any other functionality specific to the bespoke parachain's requirements are defined. The `ConstructRuntime` macro composes the overall runtime by combining FRAME modules, like the system module and custom parachain modules, to form a cohesive blockchain runtime.

This example is highly simplified and has yet to cover all aspects or complexities of a real parachain implementation. In practice, as we will see later, implementing collator logic involves more detailed considerations, including networking, consensus, state management, and interaction with the Polkadot relay chain.

How is the state machine defined

In Substrate-based blockchains, the state machine is defined implicitly through the combination of runtime modules. Each runtime module contributes to the overall state machine by defining the logic for state transitions and specifying how the blockchain's state can change over time. The state machine is a result of the interactions between these modules.

Let's revisit the example provided earlier and highlight where the state machine is implicitly defined:

```

1 // ...
2
3 // Define our parachain's runtime modules
4 pub mod my_parachain {
5     use frame_system::Config;
6     use frame_support::dispatch;
7
8     // Define the configuration for our runtime
9     pub trait Config: frame_system::Config {}
10
11     // Define our parachain's custom module
12     pub mod custom_module {
13         use super::*;

```

```

14     use frame_support::{decl_module, decl_storage, dispatch::
DispatchResult};
15
16     // Define our module's storage items
17     pub trait Config: my_parachain::Config {}
18     decl_storage! {
19         trait Store for Module<T: Config> as MyParachainModule {
20             // Define our storage items here
21             ExampleValue: u32;
22         }
23     }
24
25     // Define our module's dispatchable functions
26     decl_module! {
27         pub struct Module<T: Config> for enum Call where origin: T::
Origin {
28             // Initialize the module
29             fn deposit_event() = default;
30
31             // Example function that updates storage
32             #[weight = 10_000]
33             pub fn update_value(origin, new_value: u32) -> dispatch
::DispatchResult {
34                 let _sender = ensure_signed(origin)?;
35                 <ExampleValue<T>::put(new_value);
36                 Ok(())
37             }
38         }
39     }
40 }
41
42 // ...
43
44 // Implement collator-specific logic (block production, communication
with relay chain, etc.)
45 // ...
46
47 // ...
48
49
50

```

Listing 2: State Machine

In the myparachain module, the custommodule defines the logic for state transitions. The declstorage macro is used to declare storage items (state variables), and the declmodule macro defines dispatchable functions that can modify the state. For example, the update-value function modifies the ExampleValue storage item.

The overall state machine is a combination of the state defined in the framesystem module, the pallettimestamp module, and any custom modules like custommodule within myparachain.

The state transitions are determined by the logic within these modules, and they collectively form the state machine for the parachain. In Substrate, the state machine is dynamic and can evolve over time as runtime modules are added, modified, or upgraded. This flexibility allows developers to customise the blockchain's behavior and

state transitions according to their specific requirements.

In fact Substrate’s modular design allows for the inclusion of multiple runtime modules, each contributing to defining the state machine of a bespoke chain. These modules encapsulate specific functionalities and collectively contribute to the overall state transitions and logic of our blockchain.

In fact, designing a blockchain with multiple modules is a common and recommended practice in Substrate development. Each module can handle a distinct aspect of a blockchain’s functionality, such as managing assets, governance, identity, or any other specific feature you want to incorporate.

A parachain might have different modules covering the following areas:

- Asset Management*: This module could define state and functions related to token issuance, transfers, and balances.
- Multi-Sig: A separate module could manage user identities, access control, or any identity-related functionalities.

Each of these modules defines its own storage items, dispatchable functions, and associated logic. They work together to form the state machine, collectively governing how the blockchain’s state can change based on the actions users take through the dispatchable functions defined within these modules.

By breaking down functionalities into separate modules, it becomes easier to manage and upgrade specific parts of our blockchain without affecting the entire system. This modular approach also enhances readability, maintainability, and extensibility of our blockchain codebase.

Supporting the PinkScorpion protocol

To illustrate managing forward and reverse cryptographic transformations with a strict order of processing in a pipeline within a Substrate-based blockchain, let’s extend the outline of the previous example using multiple modules representing these transformations.

We’ll consider a simplified scenario where each transformation step is represented by a separate module:

```

1 // Forward Transformation Steps
2 pub mod forward_transform_step_one {
3     use frame_support::{decl_module, decl_storage};
4
5     pub trait Config: frame_system::Config {}
6
7     decl_storage! {
8         trait Store for Module<T: Config> as ForwardTransformStepOne {
9             // Define storage items for forward transformation step one

```

As a first phase we’ll adjust the filesystem to report protocol steps to Kafka instead of file operations. Then the Kafka consumer can be amended to make the bespoke blockchain state transitions. This allows us to evolve the performance before directly coupling into the filesystem logic.

This example now needs to be very specific to PinkScorpion and requires more thought and detail

```

10         StepOne: u32;
11     }
12 }
13
14 decl_module! {
15     pub struct Module<T: Config> for enum Call where origin: T::
Origin {
16         // Functions for executing forward transformation StepOne
17         // ...
18     }
19 }
20 }
21
22 pub mod forward_transform_step_two {
23     use frame_support::{decl_module, decl_storage};
24
25     pub trait Config: frame_system::Config {}
26
27     decl_storage! {
28         trait Store for Module<T: Config> as ForwardTransformStepTwo {
29             // Define storage items for forward transformation step two
30             StepTwo: u32;
31         }
32     }
33
34     decl_module! {
35         pub struct Module<T: Config> for enum Call where origin: T::
Origin {
36             // Functions for executing forward transformation StepTwo
37             // ...
38         }
39     }
40 }
41
42 pub mod forward_transform_step_three {
43     use frame_support::{decl_module, decl_storage};
44
45     pub trait Config: frame_system::Config {}
46
47     decl_storage! {
48         trait Store for Module<T: Config> as ForwardTransformStepThree {
49             // Define storage items for forward transformation step
three
50             StepThree: u32;
51         }
52     }
53
54     decl_module! {
55         pub struct Module<T: Config> for enum Call where origin: T::
Origin {
56             // Functions for executing forward transformation StepThree
57             // ...
58         }
59     }
60 }
61
62 // Reverse Transformation Steps (similar structure)
63
64 // Main Runtime Configuration
65 parameter_types! {
66     pub const BlockPeriod: u64 = 6;

```

```

67 }
68
69 impl forward_transform_step_one::Config for Runtime {}
70 impl forward_transform_step_two::Config for Runtime {}
71 impl forward_transform_step_three::Config for Runtime {}
72 // ... Implement other configurations for reverse transformations
73
74 construct_runtime!(
75     pub enum Runtime where
76         Block = Block,
77         NodeBlock = opaque::Block,
78         UncheckedExtrinsic = UncheckedExtrinsic
79     {
80         System: frame_system::{Module, Call, Config},
81         ForwardTransformStepOneModule: forward_transform_step_one::{
82             Module, Call, Storage, Config<T>},
83         ForwardTransformStepTwoModule: forward_transform_step_two::{
84             Module, Call, Storage, Config<T>},
85         ForwardTransformStepThreeModule: forward_transform_step_three::{
86             Module, Call, Storage, Config<T>},
87         // ... Include other reverse transformation modules
88     }
89 );
90
91 // Blockchain Storage Items
92 impl frame_system::Config for Runtime {
93     type BaseCallFilter = ();
94     type Origin = Origin;
95     type Index = u64;
96     type BlockNumber = u64;
97     type Hash = sp_core::H256;
98     type Hashing = sp_runtime::traits::BlakeTwo256;
99     type AccountId = u64;
100     type Lookup = sp_runtime::traits::IdentityLookup<Self::AccountId>;
101     type Header = sp_runtime::testing::Header;
102     type Event = ();
103     type BlockHashCount = ();
104     type DbWeight = ();
105     type BlockWeights = ();
106     type BlockLength = ();
107     type Version = ();
108     type PalletInfo = ();
109     type AccountData = ();
110     type OnNewAccount = ();
111     type OnKilledAccount = ();
112 }

```

Listing 3: State Machine

In the context of a bespoke blockchain defined with a three-stage state machine (like the example provided above), validators via the relay chain are primarily validating two critical aspects:

- **Transaction Validity:** Validators validate the transactions submitted to the bespoke blockchain. For each transformation step (StepOne, StepTwo, StepThree), validators ensure that incoming transactions comply with the predefined rules, permissions, and constraints associated with each step. For instance, validators verify that a transaction intended for StepOne meets the criteria defined for

StepOne execution. This includes validating sender authorisation, ensuring the transaction data conforms to StepOne's requirements, and verifying that the transaction doesn't violate any rules associated with that particular step.

- **Consistency and Order of State Transitions:** Validators also validate the consistency and order of state transitions within the three-stage state machine. They ensure that the transitions from one state to another (e.g., from StepOne to StepTwo and then to StepThree) follow the specified sequence and adhere to the correct progression defined by the bespoke blockchain's logic. Validators verify that the state transitions occur in the correct order, preventing out-of-sequence or unauthorised state changes. They ensure that the blockchain's state progresses through the expected transformation stages as dictated by the state machine's design.

Overall, validators play a crucial role in ensuring the correctness, integrity, and adherence to the rules specified within the bespoke blockchain's state machine. Their validation activities focus on both transaction-level validation and ensuring the proper order and progression of state transitions as defined by the bespoke blockchain's three-stage state machine.

REFERENCES

Aleph Zero

Aleph Zero is a privacy-enhancing, scalable, and decentralised blockchain platform. It is known for its unique consensus protocol based on Directed Acyclic Graph (DAG) technology combined with Byzantine Fault Tolerance (BFT), which is designed to offer fast transaction speeds, scalability, and enhanced security. Key Features of Aleph Zero:

- **Consensus Mechanism:** Aleph Zero uses a novel consensus algorithm that combines elements of DAG and BFT. This allows for high throughput and low latency in transaction processing, making it suitable for both financial transactions and complex applications like supply chain management.
- **Privacy:** One of the standout features of Aleph Zero is its commitment to privacy. It incorporates various cryptographic techniques, including zero-knowledge proofs, to ensure that transactions can be verified without revealing any sensitive information about the parties involved.

This section added as it is likely that Aleph is the means by which the protocol can be made to perform.

Understand opportunity the unique consensus protocol affords

- Scalability: The use of DAG technology allows Aleph Zero to process transactions in parallel, significantly increasing its scalability compared to traditional blockchain systems that process transactions sequentially.
- Interoperability: Aleph Zero is designed to be interoperable with other blockchains, facilitating the seamless exchange of information and value across different networks.
- Developer Friendly: The platform supports the development of decentralised applications (DApps) with a focus on security and privacy. It provides tools and frameworks that help developers build and deploy their applications efficiently.

Aleph Zero aimed to address some of the common challenges faced by traditional blockchains, such as scalability issues, high transaction fees, and privacy concerns, making it a promising technology in the blockchain space and one embraced by Lockular.

Aleph Zero is not natively built on Polkadot's Substrate, interoperability is achieved via a bridge.

Architecture: PinkScorpion protocol controlled by a blockchain

The rest of the document discusses a scheme to create a bespoke blockchain using Polkadot Substrate that integrates with a custom NFS Server. This server employs Shamir's Secret Sharing to split file information into shares for secure storage. The key operations of disassembling and reassembling file shares using Shamir's method are modeled as state transitions within a Polkadot substrate blockchain. The design aims to provide an immutable audit trail of access to the filesystem.

To address potential performance bottlenecks due to the computational intensity of Shamir's operations and the interaction with the blockchain, we propose using an off-chain aggregation approach. Here, as the filesystem applies Shamir's method, it will write these operations to an off-chain queue managed by a dedicated filesystem thread. This allows the blockchain to process these state transition transactions asynchronously, at its own pace, without being directly slowed down by real-time data processing demands.

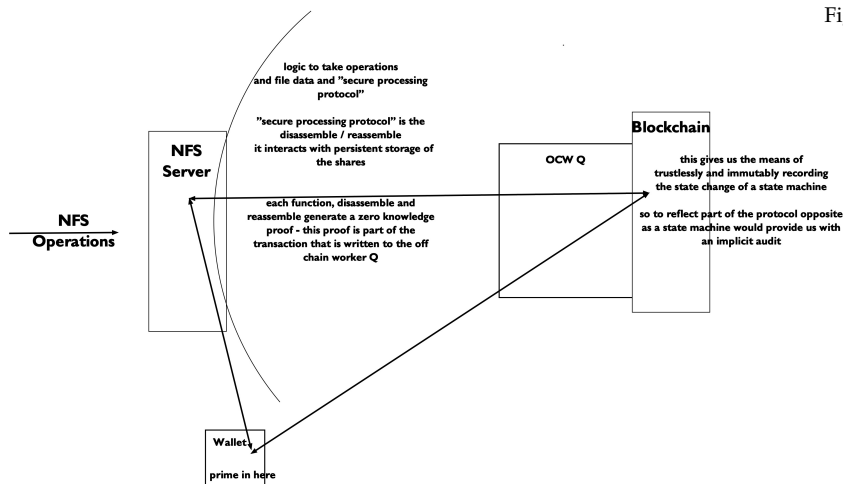


Figure 1: Overview

This setup aims to balance the security and immutability benefits of blockchain with the performance requirements of a high-load filesystem.

Our approach of using off-chain aggregation to handle the computationally intensive Shamir's Secret Sharing operations while maintaining an immutable audit trail on the blockchain is strategically sound for several reasons:

- Scalability: By offloading the Shamir operations to an off-chain

queue, we reduce the load on the blockchain, allowing it to handle more transactions and operate more efficiently. This is crucial for maintaining high performance as the system scales.

- **Security and Immutability:** The blockchain still records all critical operations, providing an immutable audit trail of file access and share reassembly. This meets the security requirements for sensitive data handling and compliance needs.
- **Asynchronous Processing:** Handling Shamir operations off-chain and allowing the blockchain to process these operations asynchronously can significantly improve the overall responsiveness and user experience of the system.
- **Flexibility:** This architecture provides flexibility in managing the balance between immediate data processing needs and the slower, immutable record-keeping process of the blockchain.

However, there are a few considerations and potential challenges, mainly due to the general case of off-chain processing:

- **Complexity:** The integration between off-chain and on-chain components adds complexity to the system, which might increase the potential for bugs and require more robust testing and maintenance strategies.
- **Consistency and State Management:** Ensuring that the off-chain state (Shamir shares) and the on-chain state (audit logs) remain consistent can be challenging. Mechanisms must be in place to handle discrepancies and failures.
- **Security of Off-Chain Components:** While the blockchain itself is secure, the off-chain components must also be secured, as they handle sensitive operations and data.
- **Latency and Timing Issues:** Depending on how the off-chain queue is managed, there might be latency or timing issues in how quickly state transitions are reflected on the blockchain, which could affect the timeliness of the audit trail.
- **Overall,** if these challenges are addressed, this approach could effectively leverage the strengths of both off-chain processing and blockchain technology to create a robust, scalable, and secure system.

That said, in the Polkadot ecosystem, you can integrate an off-chain queue more closely with the blockchain process using Substrate's off-chain features. Specifically, Substrate provides a framework for off-chain workers (OCWs) that can perform computations and manage data off the main blockchain but still within the

blockchain node's environment. This allows for a more seamless integration between on-chain and off-chain processes, mitigating many of the issues cited above.

How Off-Chain Workers Can Help

Off-chain Workers (OCWs) in Substrate run asynchronously and independently of the block production process. They can be triggered after each block is finalized and can communicate with external services, perform heavy computations, or manage queues. Here's how you can leverage OCWs for our needs:

- **Off-Chain Storage:** OCWs can utilize Substrate's off-chain storage to temporarily store data before it is committed to the blockchain. This storage is local to each node and is not consensus-critical, making it a good candidate for our off-chain queue.
- **Data Aggregation:** OCWs can aggregate data from the off-chain queue and prepare it for submission to the blockchain. This can include batching operations or performing preliminary computations necessary before updating the on-chain state. We should consider if we batch to save gas fees.
- **Scheduled Tasks:** OCWs can be scheduled to run at specific intervals or in response to specific events, allowing for regular processing of the queue.
- **Secure Data Handling:** While OCWs operate off-chain, they still maintain a high level of security and can utilise cryptographic functions provided by Substrate to ensure data integrity and security.

Implementing an Off-Chain Queue with OCWs

Here's a basic outline of how we will implement this:

- **Modify the NFS Server:** Adjust our NFS server to write operations related to Shamir's Secret Sharing into the off-chain storage managed by OCWs.
- **Develop Off-Chain Workers:** Implement OCWs that periodically read from this off-chain storage, process the data as needed, and prepare transactions to update the on-chain state.
- **Submit Transactions:** After processing, OCWs can sign and submit transactions to the blockchain to record the final state or results of operations.

- **Error Handling and Retry Logic:** Implement robust error handling and retry mechanisms in OCWs to manage failures in data processing or transaction submission.

This integration allows us to maintain the benefits of blockchain while effectively managing performance-intensive operations, keeping the system scalable and efficient.

Using Off-Chain Workers (OCWs) in Substrate does address many aspects of the security concerns associated with off-chain components, but it's important to understand the scope and limitations of this security enhancement.

Security Benefits of OCWs

- **Isolation and Encapsulation:** OCWs operate within the blockchain node's environment but are isolated from the core consensus-critical processes. This encapsulation helps in minimizing the risk of off-chain activities affecting the blockchain's integrity directly.
- **Access to Cryptographic Functions:** OCWs have access to Substrate's cryptographic functions, allowing them to securely sign transactions, verify signatures, and handle sensitive data securely. This is crucial for maintaining the authenticity and non-repudiation of the off-chain data before it's submitted to the blockchain.
- **Controlled Interaction with Blockchain:** OCWs can only interact with the blockchain through well-defined interfaces, specifically through submitting signed transactions. This controlled interaction reduces the risk of unauthorized or malicious modifications to the blockchain state.

Limitations and Remaining Security Concerns

- **Node-Level Security:** Since OCWs run on individual nodes, the security of each node becomes crucial. If a node is compromised, the OCWs running on that node could potentially be manipulated. Ensuring the security of the hardware and operating environment of nodes is essential.
- **Data Transmission Security:** While OCWs can securely handle data at the node, the transmission of data between nodes or from external sources to nodes needs to be secured separately, typically using secure communication protocols like TLS.
- **Off-Chain Data Storage:** Data stored off-chain, even temporarily, is not protected by the blockchain's consensus mechanisms. This

data could be at risk if not properly encrypted or if the storage medium is not secure.

- **Dependency on External Data and Services:** If OCWs rely on external data or services, the security and reliability of these external components can impact the overall system. For instance, data feeds or APIs that OCWs interact with need to be secure and trustworthy. In our case I believe we are ok here.

Mitigation Strategies

- **Node Security:** We should implement robust security measures for the nodes, including secure boot mechanisms, operating system hardening, and intrusion detection systems.
- **Data Encryption:** Encrypt sensitive data handled by OCWs both in transit and at rest, using strong encryption protocols.
- **Regular Audits and Updates:** Regularly audit the code and security infrastructure of OCWs and update them to patch any vulnerabilities.
- **Redundancy and Consensus:** For critical operations, consider having multiple OCWs across different nodes perform the same computations and use a consensus mechanism among them before submitting transactions to the blockchain.

In summary, while OCWs enhance the security of off-chain components by leveraging the blockchain node's capabilities, they do not eliminate all security risks. A comprehensive security strategy that addresses both on-chain and off-chain components is essential for maintaining the overall security of the system.

Integrity of Blockchain nodes

Given the scenario where users access the NFS Server using a Polkadot.js wallet containing a Shamir prime for decrypting their files, ensuring the integrity of the blockchain nodes used is crucial. A streamlined approach to achieve this is outlined here, focusing on user registration and session validation through the blockchain:

- **User Registration on the Blockchain:**
 - **User Wallet:** Each user registers on the blockchain via their Polkadot.js wallet. This registration can include their public key and any other relevant identifiers.

- Smart Contract: Implement a smart contract on the blockchain that handles user registrations. This contract stores user details securely and can be queried to verify user identities. So this contract is the general RBAC contract.
- Session Initialization:
 - Session Token: When a user accesses the NFS Server, they initiate a session by creating a session token or similar credential, which is then signed using their private key (from their wallet).
 - Blockchain Verification: This token is sent to the blockchain (via a transaction or a direct smart contract call), where it's verified against the registered user details.
- Filesystem Verification:
 - Smart Contract Query: Before allowing access to any files, the NFS Server queries the blockchain to verify the session token against the active user sessions registered in the smart contract.
 - Access Control: If the session token is valid and matches the registered details on the blockchain, the NFS Server grants access to the user's files. If not, access is denied.
- Blockchain Node Integrity:
 - Node Verification: To ensure the NFS Server is interacting with the correct blockchain nodes, use a list of trusted nodes maintained either in a secure configuration file or through a decentralized registry on the blockchain itself.
 - TLS and Certificates: Secure communications between the NFS Server and the blockchain nodes using TLS, where each node presents a certificate that the NFS Server can verify against a list of trusted certificates.

This approach leverages blockchain for user authentication and session management, ensuring that the NFS Server interacts only with verified users and trusted blockchain nodes, thereby enhancing the overall security and integrity of the system.

Attesting to the Shamir steps

ZKPs can be used to validate transactions on the blockchain without revealing the underlying data that generated these transactions. A proof could be generated to confirm that a file operation was correctly performed according to the Shamir's Secret Sharing protocol, without revealing the content of the files or the shares themselves.

Introducing zero-knowledge proofs (ZKPs) to verify file operations performed according to Shamir's Secret Sharing protocol without revealing the content of the files or the shares themselves involves several steps. Here's how we can integrate this into our architecture:

- Define the ZKP Scheme First, you need to define the specific ZKP scheme that will be used to verify the operations. This involves:
 - Choosing the right type of ZKP: Depending on our needs, you might choose zk-SNARKs, zk-STARKs, or another variant. zk-SNARKs are widely used due to their efficiency in verification times but require a trusted setup. zk-STARKs, on the other hand, do not require a trusted setup and are post-quantum secure but are generally larger in size.
 - Creating the circuit: Define a computational circuit that represents the logic of Shamir's Secret Sharing operations. This circuit will be used to generate and verify proofs. A circuit is required for disassembly and a circuit for reassembly.
- Integration Points
 - During File Operations: Integrate the generation of ZKPs into the process where files are split into shares or reassembled. This can be done on the client side (in the user's environment) or by the off-chain workers if they handle such operations.
 - Verification by Smart Contract: Deploy a smart contract capable of verifying the proofs submitted along with transactions that record file operations on the blockchain. This contract checks the validity of the operations without needing to see the underlying data.
- Implementation Steps
 - Proof Generation: When a file is split or shares are combined, the corresponding ZKP is generated. This proof attests that the operation was performed correctly according to the rules defined in the computational circuit.
 - Submitting Proofs to Blockchain: Along with recording the operation on the blockchain, the proof is also submitted. This can be part of the transaction data.
 - Smart Contract for Proof Verification: The smart contract will use the ZKP verification algorithm to ensure the proof is valid. If the proof is valid, the operation is recorded as successful; otherwise, it is rejected.
- Security and Privacy Considerations

- Confidentiality: Ensure that no sensitive data is leaked during proof generation or verification. The ZKP should only prove the correctness of the operation, not reveal any data.
- Efficiency: Consider the computational and storage overhead introduced by ZKPs. Optimize the ZKP setup to balance security, privacy, and performance.

Architecture Summary

Core components comprise:

- NFS Server: Manages secure file storage and access, utilising Shamir's Secret Sharing to split file information into shares.
- Polkadot.js Wallet: Used by users to store their unique Shamir prime and interact with the blockchain.
- Bespoke Blockchain (Built on Polkadot Substrate): Custom blockchain that models Shamir's Secret Sharing operations as state transitions. This blockchain uses smart contracts for user registration and session management, and integrates off-chain workers for efficient processing.
- Smart Contracts and Off-Chain Workers: Smart contracts handle user registration and session management. Off-chain workers process Shamir operations off-chain to enhance performance and then queue these operations for asynchronous processing on the blockchain.
- Zero-Knowledge Proofs (ZKPs): Used to verify the correctness of Shamir's Secret Sharing operations (disassembly and reassembly) without revealing the content of the files or the shares themselves.

Workflow steps:

- User Registration:
 - Users register on the blockchain via their Polkadot.js wallet, providing necessary credentials and their public key.
 - Registration details are securely stored in a smart contract on the blockchain.
- Session Management:
 - Users initiate sessions by creating a session token signed with their private key when accessing the NFS Server.
 - This token is verified against registered details on the blockchain via a smart contract.

- File Access and Shamir Operations:
 - Upon successful session verification, users can access and operate on their files.
 - File operations involve Shamir's Secret Sharing, where file data is split or combined using the user's part of the key from their wallet.
 - Each operation (splitting or combining shares) corresponds to a state transition in the bespoke blockchain, recorded and managed by the blockchain's state machine.
- Zero-Knowledge Proof Generation and Verification:
 - During file operations, ZKPs are generated to prove that operations were performed correctly according to Shamir's Secret Sharing protocol.
 - Proofs are submitted along with the state transition data to the blockchain via OCWs.
 - Smart contracts on the blockchain verify these proofs to ensure the integrity and correctness of the operations without accessing the actual data.
- Off-Chain Workers:
 - Off-chain workers process Shamir operations and ZKP generation off-chain to enhance performance.
 - These operations, along with their proofs, are then queued and asynchronously processed on the blockchain, ensuring efficient and scalable handling of data-intensive tasks.
- Blockchain Node Integrity and Security:
 - The NFS Server communicates only with verified blockchain nodes, ensuring node integrity through a list of approved nodes or certificates.
 - All communications are secured using TLS.

Security Features include:

- Immutable Audit Trails: The blockchain provides an immutable record of all Shamir operations as state transitions, along with ZKP verifications, enhancing transparency and security.
- Decentralized Trust: Utilizes blockchain for decentralized security, reducing the risk of tampering and reliance on a single point of failure.

- Encrypted Communications: Uses TLS for secure communication between the NFS Server and blockchain nodes.
- Role-Based Access and Multi-Factor Authentication: Ensures that only authorized users can access their files and make configuration changes.

This architecture effectively integrates the security and immutability of blockchain technology with the computational efficiency provided by off-chain workers and the privacy-preserving capabilities of zero-knowledge proofs, specifically tailored to manage the complexities of Shamir's Secret Sharing in a file storage environment.

References